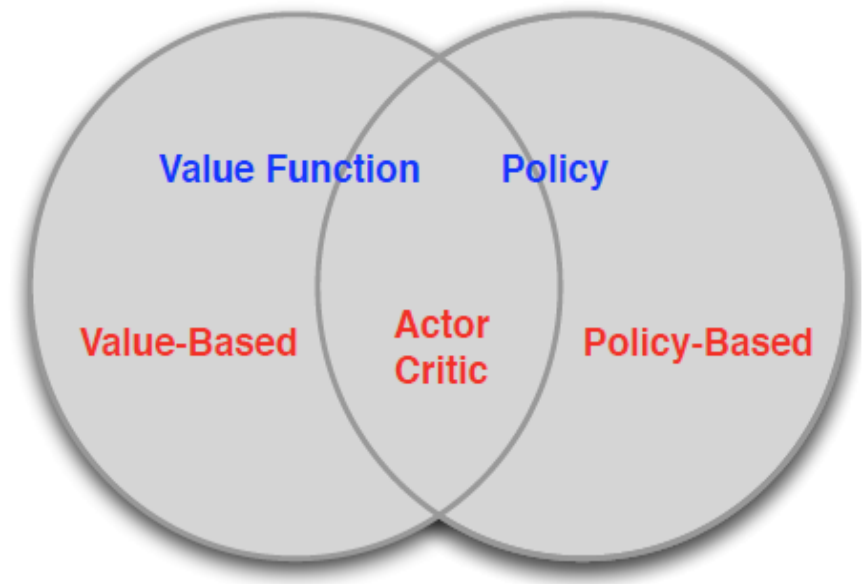


Value-Based and Policy-Based RL

- Value Based
 - Learnt Value Function
 - Implicit policy (e.g. ϵ -greedy)
- Policy Based
 - No Value Function
 - Learnt Policy
- Actor-Critic
 - Learnt Value Function
 - Learnt Policy



References

- A3C
 - Asynchronous Methods for Deep Reinforcement Learning
 - Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mirza, Alex Graves, Timothy P. Lillicrap, Tim Harley, David Silver, Koray Kavukcuoglu
 - Google DeepMind, Montreal Institute for Learning Algorithms (MILA), University of Montreal
- TRPO
 - Trust Region Policy Optimization
 - ▶ Schulman, J., et al. Trust region policy optimization. In: *International Conference on Machine Learning*. 2015. p. 1889-1897.
- PPO
 - Proximal Policy Optimization Algorithms
 - ▶ Schulman, J., et al. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- Contributors for the slides include: 蔡承倫, 林九州, 何國豪, etc.

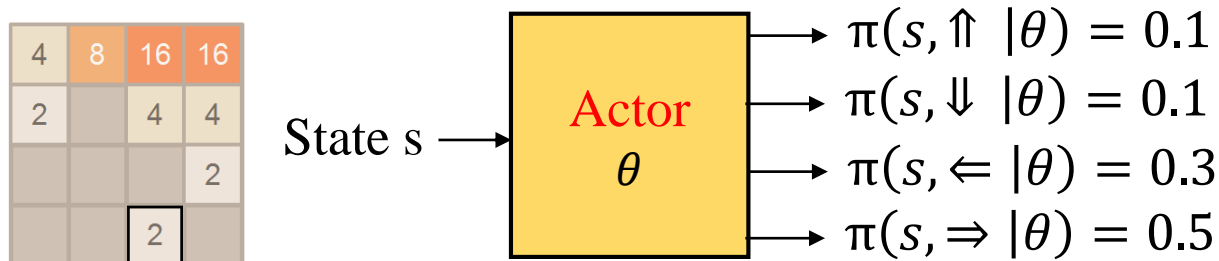

Policy-Based Reinforcement Learning

- Policy Gradient
- Actor-Critic (Discrete actions)
- A3C (Asynchronous Advantage Actor-Critic)
- TRPO & PPO
- DDPG (Deep Deterministic Policy Gradient)
 - ▶ TD3
 - ▶ SAC



An Example

- DQN outputs the values of actions. (Up/Down/Left/Right)
- Actor outputs the policy, probability of selecting actions.



Advantages of Policy-Based RL

- Advantages:

- Better convergence properties
 - ▶ Recall grid world with equal policy for left/up/right/down operations.
- Effective in high-dimensional or continuous action spaces
- Can learn stochastic policies

- Disadvantages:

- Typically converge to a local rather than global optimum
- Evaluating a policy is typically inefficient and high variance

Policy Objective Functions

- Goal:
 - given policy $\pi_\theta(s, a)$ with parameters θ , **find best θ**
 - ▶ What does the best mean?
 - ▶ How do we **measure the quality of a policy π_θ** ?
- In episodic environments we can use the **start value**
- In continuing environments we can use the **average value**

$$J_0(\theta) = V^{\pi_\theta}(s_0) = \mathbb{E}_{\pi_\theta}[v_0]$$

$$J_{avV}(\theta) = \sum_s d^{\pi_\theta}(s) V^{\pi_\theta}(s)$$

- Or the average reward per time-step

$$J_{avR}(\theta) = \sum_s d^{\pi_\theta}(s) \sum_a \pi_\theta(s, a) R_s^a$$

- Where $d^{\pi_\theta}(s)$ is stationary distribution of Markov chain for π_θ



Policy Optimization

- Policy based reinforcement learning is an **optimization** problem
 - Find θ that maximizes $J(\theta)$
- Some approaches do not use gradient
 - Hill climbing
 - Simplex / amoeba / Nelder Mead
 - Genetic algorithms
- Greater efficiency often possible using gradient
 - Gradient descent
 - Conjugate gradient
 - Quasi-newton
- We focus
 - on gradient descent, many extensions possible
 - And on methods that exploit sequential structure



Policy Gradient


- Let $J(\theta)$ be any policy objective function
- Policy gradient algorithms search for a local maximum in $J(\theta)$ by ascending the gradient of the policy, w.r.t. parameters θ

$$\Delta\theta = \alpha \nabla_{\theta} J(\theta)$$

- Where $\nabla_{\theta} J(\theta)$ is the policy gradient

$$\nabla_{\theta} J(\theta) = \begin{pmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_n} \end{pmatrix}$$

- and α is a step-size parameter



Computing Gradients By Finite Differences

- To evaluate policy gradient of $\pi_{\theta}(s, a)$
- For each dimension $k \in [1, n]$
 - Estimate k th partial derivative of objective function w.r.t. θ
 - By perturbing θ by small amount ϵ in k th dimension
$$\frac{\partial J(\theta)}{\partial \theta_k} \approx \frac{J(\theta + \epsilon u_k) - J(\theta)}{\epsilon}$$
 - ▶ where u_k is unit vector with 1 in k th component, 0 elsewhere
 - Uses n evaluations to compute policy gradient in n dimensions
- Simple, noisy, inefficient - but sometimes effective
- Works for arbitrary policies, even if policy is not differentiable

Policy Gradient (One Step)

- Consider a simple class of **one-step MDPs**
- Starting in state $s_0 \sim d(s)$
- Terminating after one time-step with reward $r = R_{s,a}$
- Use likelihood ratios to compute the policy gradient

$$\begin{aligned}
 J_0(\theta) &= V^{\pi_\theta}(s_0) = \mathbb{E}_{\pi_\theta}[r] = \sum_{a \in A} \pi_\theta(s_0, a) R_{s_0, a} \\
 \nabla_\theta J_0(\theta) &= \sum_{a \in A} \nabla_\theta \pi_\theta(s_0, a) R_{s_0, a} \\
 &= \sum_{a \in A} \pi_\theta(s_0, a) \nabla_\theta \log \pi_\theta(s_0, a) R_{s_0, a} \\
 &= \mathbb{E}_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) \cdot r] \quad \text{Score function}
 \end{aligned}$$

Let $s_0 \sim d(s)$

$$\begin{aligned}
 J(\theta) &= \mathbb{E}_{d(s), \pi_\theta}[r] \\
 &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in A} \pi_\theta(s, a) R_{s, a} \\
 \nabla_\theta J(\theta) &= \mathbb{E}_{d(s), \pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) \cdot r]
 \end{aligned}$$



Score Function

- We now compute the policy gradient analytically
- Assume
 - policy π_θ is differentiable whenever it is non-zero
 - we know the gradient $\nabla_\theta \pi_\theta(s, a)$

- **Likelihood** ratios exploit the following identity

$$\begin{aligned}\nabla_\theta \pi_\theta(s, a) &= \pi_\theta(s, a) \frac{\nabla_\theta \pi_\theta(s, a)}{\pi_\theta(s, a)} \\ &= \pi_\theta(s, a) \nabla_\theta \log \pi_\theta(s, a)\end{aligned}$$

- $\nabla_\theta \log \pi_\theta(s, a)$ is called the **score function**.

Softmax Policy

- Probability of action is proportional to exponentiated weight

$$\pi_{\theta}(s, a) \propto e^{\phi(s, a)^T \theta}$$

- Weight actions using linear combination of features $\phi(s, a)^T \theta$

- The score function is

$$\nabla_{\theta} \log \pi_{\theta}(s, a) = \phi(s, a) - \mathbb{E}_{\pi_{\theta}}[\phi(s, \cdot)]$$

- Example:

- In Computer Go, Silver used this to solve a problem
 - ▶ Simulation Balancing



Gaussian Policy

- In continuous action spaces, a Gaussian policy is natural
- Mean is a linear combination of state features $\mu_{\theta}(s) = \phi(s)^T \theta$
- Variance may be fixed σ^2 or can also be parametrized
- Policy is Gaussian, $a \sim \mathcal{N}(\mu_{\theta}(s), \sigma^2)$
- The score function is

$$\nabla_{\theta} \log \pi_{\theta}(s, a) = \frac{(a - \mu_{\theta}(s)) \phi(s)}{\sigma^2}$$

Score Function Gradient Estimator

- Consider an expectation $\mathbb{E}_{x \sim p(x|\theta)}[f(x)]$.

- The gradient w.r.t. θ is:

$$\nabla_{\theta} \mathbb{E}_x[f(x)] = \mathbb{E}_x[f(x) \nabla_{\theta} \log p(x|\theta)]$$

- Just sample $x_i \sim p(x|\theta)$, and compute
$$\hat{g}_i = f(x_i) \nabla_{\theta} \log p(x_i|\theta)$$
- Need to be able to compute and differentiate density $p(x|\theta)$ w.r.t. θ
- This gives us an unbiased gradient estimator.
- Note: $\pi_{\theta}(s, a)$ can be viewed as $p(x|\theta)$.



One-Step MDPs


- Consider a simple class of **one-step MDPs**
- Starting in state $s \sim d(s)$
- Terminating after one time-step with reward $r = R_{s,a}$
- Use likelihood ratios to compute the policy gradient

$$\begin{aligned} J(\theta) &= \mathbb{E}_{\pi_{\theta}}[r] \\ &= \sum_{s \in S} d(s) \sum_{a \in A} \pi_{\theta}(s, a) R_{s,a} \\ \nabla_{\theta} J(\theta) &= \sum_{s \in S} d(s) \sum_{a \in A} \pi_{\theta}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a) R_{s,a} \\ &= \mathbb{E}_{\pi_{\theta}}[\nabla_{\theta} \log \pi_{\theta}(s, a) \cdot r] \end{aligned}$$



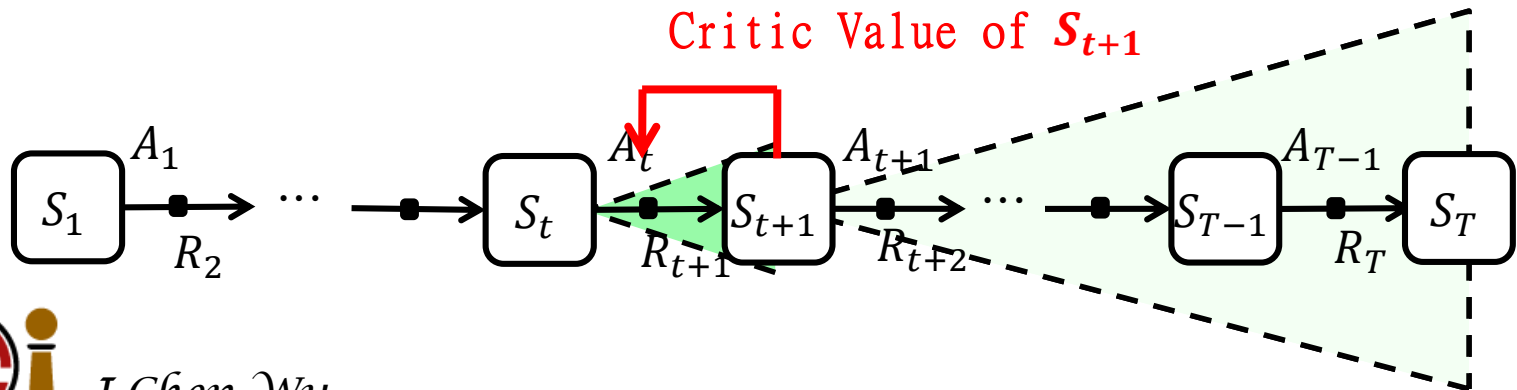
Problem of REINFORCE

- Problem: Monte-Carlo policy gradient still has **high variance**



- Solution: Actor Critic

- Policy gradient based on the Critic value of S_{t+1}



Policy-Based Reinforcement Learning

- Policy Gradient
- Actor-Critic (Discrete actions)
- A3C (Asynchronous Advantage Actor-Critic)
- TRPO & PPO
- DDPG (Deep Deterministic Policy Gradient)
 - ▶ TD3
 - ▶ SAC



Reducing Variance Using a Critic

- We use a **critic** to estimate the action-value function,

$$Q_w(s_t, a_t) \approx Q^{\pi_\theta}(s, a)$$

- Actor-critic algorithms maintain **two** sets of parameters

- **Critic**: Updates action-value function parameters w
- **Actor**: Updates policy parameters θ , in direction suggested by critic

- Actor-critic algorithms follow **an approximate policy gradient**

$$\nabla_\theta J(\theta) \approx \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) \cdot Q_w(s, a)]$$


$$\Delta\theta = \alpha \nabla_\theta \log \pi_\theta(s, a) \cdot Q_w(s, a)$$



Estimating the Action-Value Function

- The **critic** is solving a familiar problem: **policy evaluation**
- But, how good is policy π_θ for current parameters θ ?
- This problem was explored in previous two chapters, e.g.
 - Monte-Carlo policy evaluation
 - Temporal-Difference learning
 - TD(λ)
- Could also use e.g. least-squares policy evaluation

Actor-Critic (Discrete Action Space)



- Use two networks: an **actor** and a **critic**

- **Critic** estimates the action-value function

- ▶ Gradient:



$$\nabla_{\omega} L_Q(s_t, a_t | \omega) = ((r_{t+1} + \gamma Q(s_{t+1}, a' | \omega)) - Q(s_t, a_t | \omega)) \nabla_{\omega} Q(s_t, a_t | \omega)$$

- **Actor** updates policy in direction suggested by critic

- ▶ Gradient (approximate policy gradient):

$$J(\theta) = E_{s,a}^{\pi_{\theta}} [Q(s, a | \omega)]$$

$$\nabla_{\theta} J(\theta) = E_{s,a}^{\pi_{\theta}} [\nabla_{\theta} \log \pi(s_t, a_t | \theta) Q(s_t, a_t | \omega)]$$



Actor-Critic (Discrete Action Space)

- Using linear value function approx. $Q_w(s, a) = \phi(s, a)^T w$
 - **Critic**: Updates w by linear TD(0)
 - **Actor**: Updates θ by policy gradient

```
function QAC
  Initialise  $s, \theta$ 
  Sample  $a \sim \pi_\theta$ 
  for each step do
    Sample reward  $r = \mathcal{R}_s^a$ ; sample transition  $s' \sim \mathcal{P}_{s, \cdot}^a$ .
    Sample action  $a' \sim \pi_\theta(s', a')$ 
     $\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$ 
     $\theta = \theta + \alpha \nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)$ 
     $w \leftarrow w + \beta \delta \phi(s, a)$ 
     $a \leftarrow a', s \leftarrow s'$ 
  end for
end function
```



Policy-Based Reinforcement Learning

- Policy Gradient
- Actor-Critic (Discrete actions)
- A3C (Asynchronous **Advantage Actor-Critic**)
- TRPO & PPO
- DDPG (Deep Deterministic Policy Gradient)
 - ▶ TD3
 - ▶ SAC



Estimating the Advantage Function (1)

- The advantage function can significantly reduce variance of policy gradient
- So the critic should really estimate the advantage function
 - For example, by estimating both $V^{\pi_\theta}(s)$ and $Q^{\pi_\theta}(s, a)$
 - Using two function approximators and two parameter vectors,
$$V_v(s) \approx V^{\pi_\theta}(s)$$
$$Q_w(s, a) \approx Q^{\pi_\theta}(s, a)$$
$$A(s, a) = Q_w(s, a) - V_v(s)$$
- And updating both value functions by e.g. TD learning

Estimating the Advantage Function (2)

- For the true value function $V^{\pi_\theta}(s)$, the TD error δ^{π_θ}

$$\delta^{\pi_\theta} = r + \gamma V^{\pi_\theta}(s') - V^{\pi_\theta}(s)$$

- is an unbiased estimate of the advantage function

$$\begin{aligned}\mathbb{E}_{\pi_\theta}[\delta^{\pi_\theta} | s, a] &= \mathbb{E}_{\pi_\theta}[r + \gamma V^{\pi_\theta}(s') | s, a] - V^{\pi_\theta}(s) \\ &= Q^{\pi_\theta}(s, a) - V^{\pi_\theta}(s) \\ &= A^{\pi_\theta}(s, a)\end{aligned}$$

- So we can use the TD error to compute the policy gradient

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta}[\nabla_\theta \log \pi_\theta(s, a) \delta^{\pi_\theta}]$$

- In practice we can use an approximate TD error

$$\delta_v = r + \gamma V_v(s') - V_v(s)$$

- This approach only requires one set of critic parameters v



Actors at Different Time-Scales

- The policy gradient can also be estimated at many time-scales

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]$$

- Monte-Carlo policy gradient uses error from complete return

$$\Delta\theta = \alpha(v_t - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Actor-critic policy gradient uses the one-step TD error

$$\Delta\theta = \alpha(r + \gamma V_v(s_{t+1}) - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Advantage Actor-critic (A2C or A3C) policy gradient uses the $(k+1)$ -step TD error

$$\Delta\theta = \alpha(v_t^{(k)} - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Some policy gradient algorithms (like PPO) uses TD(λ) error

$$\Delta\theta = \alpha(v_t^{\lambda} - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$



Actors at Different Time-Scales

- The policy gradient can also be estimated at many time-scales

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]$$

- Monte-Carlo policy gradient uses error from complete return

$$\Delta\theta = \alpha(v_t - V_v(s_t)) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Actor-critic policy gradient uses the one-step TD error

$$\Delta\theta = \alpha(\delta_t) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Advantage Actor-critic (A2C or A3C) policy gradient uses the $(k+1)$ -step TD error $= A^{(k+1)}$

$$\Delta\theta = \alpha(\delta_t + \gamma\delta_{t+1} + \cdots + \gamma^k\delta_{t+k}) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

- Some policy gradient algorithms (like PPO) uses TD(λ) error

$$\Delta\theta = \alpha(\delta_t + \lambda\gamma\delta_{t+1} + \cdots + (\lambda\gamma)^k\delta_{t+k} + \cdots) \nabla_{\theta} \log \pi_{\theta}(s_t, a_t)$$

$= A_t^{GAE}$: Also called GAE (Generalized Advantage Estimator)



Summary of Policy Gradient Algorithms

- The **policy gradient** has many equivalent forms

$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) v_t]$	REINFORCE
$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q^w(s, a)]$	Q Actor-Critic
$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \delta]$	TD Actor-Critic
$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{(k)}]$	Advantage Actor-Critic
$= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{GAE}]$	TD(λ) Actor-Critic

- Each leads a stochastic gradient ascent algorithm



Appendix for Advantages and TD(λ) Errors

- TD errors
- n-step TD errors
- GAE
- Eligibility Trace

Appendix: n-Step TD Errors

- Sum them up, becoming n-step TD errors.

$$\begin{aligned}
 \delta_t^V &= -V(s_t) + r_t + \cancel{\gamma V(s_{t+1})} \\
 \gamma * \delta_{t+1}^V &= \gamma * (-\cancel{V(s_{t+1})} + r_{t+1} + \cancel{\gamma V(s_{t+2})}) \\
 \gamma^2 * \delta_{t+2}^V &= \gamma^2 * (-\cancel{V(s_{t+2})} + r_{t+2} + \cancel{\gamma V(s_{t+3})}) \\
 &\vdots \\
 + \quad \gamma^n * \delta_{t+n}^V &= \gamma^n * (-\cancel{V(s_{t+n})} + r_{t+n} + \cancel{\gamma V(s_{t+n+1})})
 \end{aligned}$$

$$\begin{aligned}
 &\delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V + \dots + \gamma^n \delta_{t+n}^V \\
 &= -V(s_t) + r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{n+1} V(s_{t+n+1}) \\
 &= -V(s_t) + G_t^{(n+1)} \\
 &= \hat{A}_t^{(n+1)}
 \end{aligned}$$

If $n \rightarrow \infty$, it becomes MC learning. Why?



Appendix: n-Step TD Errors

- n-step TD errors:

$$\hat{A}_t^{(1)} :=$$

$$\delta_t^V$$

$$\hat{A}_t^{(2)} :=$$

$$(\delta_t^V + \gamma \delta_{t+1}^V)$$

$$\hat{A}_t^{(3)} :=$$

$$(\delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V)$$

$$\vdots$$

$$\hat{A}_t^{(n)} :=$$

$$\sum_{k=1}^n \gamma^{k-1} * \delta_{t+k-1}^V$$



Appendix: n-Step TD Errors and GAE

- Weighted n-step TD errors:
 - The same trick as TD(λ)
- Then, sum them up.

$$\begin{aligned}
 (1 - \lambda) * \hat{A}_t^{(1)} &:= (1 - \lambda) & * \delta_t^V \\
 (1 - \lambda)\lambda * \hat{A}_t^{(2)} &:= (1 - \lambda)\lambda & * (\delta_t^V + \gamma\delta_{t+1}^V) \\
 (1 - \lambda)\lambda^2 * \hat{A}_t^{(3)} &:= (1 - \lambda)\lambda^2 & * (\delta_t^V + \gamma\delta_{t+1}^V + \gamma^2\delta_{t+2}^V)
 \end{aligned}$$

$$\vdots$$

$$+ \quad (1 - \lambda)\lambda^{n-1} * \hat{A}_t^{(n)} := (1 - \lambda) \sum_{k=1}^n \gamma^{k-1} * \delta_{t+k-1}^V$$

$$\hat{A}_t^{GAE(\gamma, \lambda)} = (1 - \lambda)(\hat{A}_t^{(1)} + \lambda\hat{A}_t^{(2)} + \lambda^2\hat{A}_t^{(3)} + \dots + \lambda^{n-1}\hat{A}_t^{(n)} + \dots)$$



Appendix: n-Step TD Errors and GAE

- The sum of exponentially-weighted TD residuals denoted as $\hat{A}_t^{GAE(\gamma, \lambda)}$ (actually equals to $G_t^\lambda - V(S_t)$ for $TD(\lambda)$)

$$\begin{aligned}
 \hat{A}_t^{GAE(\gamma, \lambda)} &= (1 - \lambda) \left(\hat{A}_t^{(1)} + \lambda \hat{A}_t^{(2)} + \lambda^2 \hat{A}_t^{(3)} + \dots + \lambda^{n-1} \hat{A}_t^{(n)} + \dots \right) \\
 &= (1 - \lambda) \left((\delta_t^V) + \lambda (\delta_t^V + \gamma \delta_{t+1}^V) + \lambda^2 (\delta_t^V + \gamma \delta_{t+1}^V + \gamma \delta_{t+2}^V) + \dots \right) \\
 &= (1 - \lambda) \left(\begin{array}{c} \delta_t^V (1 + \lambda + \lambda^2 + \dots) + \\ \gamma \lambda \delta_{t+1}^V (1 + \lambda + \lambda^2 + \dots) + \\ (\gamma \lambda)^2 \delta_{t+2}^V (1 + \lambda + \lambda^2 + \dots) + \\ \dots \end{array} \right) \\
 &= (1 - \lambda) \left(\delta_t^V \left(\frac{1}{1 - \lambda} \right) + \gamma \lambda \delta_{t+1}^V \left(\frac{1}{1 - \lambda} \right) + (\gamma \lambda)^2 \delta_{t+2}^V \left(\frac{1}{1 - \lambda} \right) + \dots \right) \\
 &= \sum_{n=0}^{\infty} (\gamma \lambda)^n \delta_{t+n}^V = \delta_t^V + \lambda \gamma \delta_{t+1}^V + \dots + (\lambda \gamma)^k \delta_{t+k}^V + \dots
 \end{aligned}$$



Appendix: Recall $TD(\lambda)$

- λ -return G_t^λ :
 - combines all n -step returns $G_t^{(n)}$
- Using weight $(1 - \lambda) \lambda^{n-1}$

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$


or (in the case of termination)

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + (1 - \lambda) \sum_{n=T-t-1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + \lambda^{T-t-1} G_t$$

- Forward-view $TD(\lambda)$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$



Appendix: GAE and Eligibility Trace

- Eligibility trace:

$$E_0(s) = 0$$

$$E_t(s) = (\gamma \lambda) E_{t-1}(s) + 1(S_t = s)$$

$$\begin{aligned} \hat{A}_t^{GAE(\gamma, \lambda)} &= 1 \delta_t^V + \lambda \gamma \delta_{t+1}^V + (\lambda \gamma)^2 \delta_{t+2}^V + \dots + (\lambda \gamma)^k \delta_{t+k}^V + \dots \\ \hat{A}_{t+1}^{GAE(\gamma, \lambda)} &= 1 \delta_{t+1}^V + \lambda \gamma \delta_{t+2}^V + \dots + (\lambda \gamma)^{k-1} \delta_{t+k}^V + \dots \\ \hat{A}_{t+2}^{GAE(\gamma, \lambda)} &= 1 \delta_{t+2}^V + \dots + (\lambda \gamma)^{k-2} \delta_{t+k}^V + \dots \\ &\dots \end{aligned}$$

$E_t(s_t)$ $E_t(s_{t+1})$ $E_t(s_{t+2})$




Policy-Based Reinforcement Learning

- Policy Gradient
- Actor-Critic (Discrete actions)
- **A3C (Asynchronous Advantage Actor-Critic)**
- TRPO & PPO
- DDPG (Deep Deterministic Policy Gradient)
 - ▶ TD3
 - ▶ SAC




Asynchronous Advantage Actor-critic (A3C)

- Asynchronous Lock-Free Reinforcement Learning
 - Use two main ideas to make the algorithm practical:
 - ▶ Multiple threads on a **single machine**
 - ▶ Multiple actor-learners applying **online updates** in parallel (**no experience replay**)



Asynchronous Advantage Actor-critic (A3C)

- Instead of experience replay, we **asynchronously** execute **multiple agents** in parallel.
 - Decorrelate the agents' data into a more stationary process
 - Enable a much larger spectrum of fundamental **on-policy RL algorithms**
- For each worker (asynchronous part):
 - Copy all parameters from the global network.
 - keep playing and computing gradients.
 - Every N iterations:
 1. **Update** all gradients to the global network.
 2. **Copy** all new parameters from the global network



Asynchronous Advantage Actor-critic (A3C)

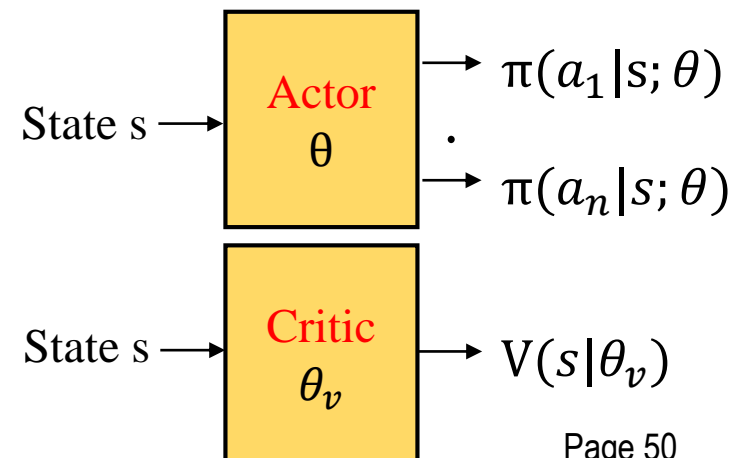
- Asynchronous advantage actor-critic(A3C) maintains a policy $\pi(a_t|s_t; \theta)$ and an estimate of the value function $V(s_t, \theta_v)$.

- The update performed by the algorithm can be seen as

$$\nabla_{\theta} \log \pi(a_t|s_t; \theta) \underbrace{A(s_t, a_t; \theta, \theta_v)}_{\text{Advantage}} = \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k}; \theta_v) - V(s_t; \theta_v)$$

– Make k -step operations, and then calculate advantages backwards.

- Intuitively, the network should be



Asynchronous Advantage Actor-critic (A3C)

- Asynchronous advantage actor-critic(A3C) maintains a policy $\pi(a_t|s_t; \theta)$ and an estimate of the value function $V(s_t, \theta_v)$.

- The update performed by the algorithm can be seen as

$$\nabla_{\theta} \log \pi(a_t | s_t; \theta) A(s_t, a_t; \theta, \theta_v) \xrightarrow{\quad} \sum_{i=0}^{k-1} \gamma^i r_{t+i} + \gamma^k V(s_{t+k}; \theta_v) - V(s_t; \theta_v)$$


- Make k -step operations, and then calculate advantages backwards.

- Typically use a convolutional neural network that has two heads:

- one **softmax output** for the policy $\pi(a_t|s_t; \theta)$

- one **output** for the value function $V(s_t; \theta_v)$

- all non-output layers are shared



Asynchronous Advantage Actor-critic (A3C)

repeat

Sync $\theta' = \theta, \theta'_v = \theta_v$

$t_{start} = t$

Get state s_t

repeat (note: $t = t + 1$)

Perform a_t according to policy $\pi(a_t|s_t; \theta')$

Receive s' and reward r

until terminal s_t or $t - t_{start} == t_{max}$

$R = \begin{cases} 0 & \text{for terminal } s' \\ V(s_t, \theta'_v) & \text{for non-terminal } s' \end{cases}$

θ, θ_v : global shared parameters

T : global shared counter

θ', θ'_v : thread specific parameters

t : thread step counter

(note: $t = t_{start} + t_{max}$,
if not terminal)

for $i \in \{t - 1, \dots, t_{start}\}$ **do**

$R \leftarrow r_i + \gamma R$

Accumulate gradients wrt θ' : $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta') (R - V(s_i; \theta'_v))$

Accumulate gradients wrt θ'_v : $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

end for

Perform asynchronous update of θ using $d\theta$ and of θ_v using $d\theta_v$.

until $T \rightarrow T_{max}$



Experiments – A3C

Method	Training Time	Mean	Median
DQN (from [Nair et al., 2015])	8 days on GPU	121.9%	47.5%
Gorila [Nair et al., 2015]	4 days, 100 machines	215.2%	71.3%
Double DQN [Van Hasselt et al., 2015]	8 days on GPU	332.9%	110.9%
Dueling Double DQN [Wang et al., 2015]	8 days on GPU	343.8%	117.1%
Prioritized DQN [Schaul et al., 2015]	8 days on GPU	463.6%	127.6%
A3C, FF	1 day on CPU	344.1%	68.2%
A3C, FF	4 days on CPU	496.8%	116.6%
A3C, LSTM	4 days on CPU	623.0%	112.6%

Table 1: Mean and median human-normalized scores on 57 Atari games using the human starts evaluation metric.



Policy-Based Reinforcement Learning

- Policy Gradient
- Actor-Critic (Discrete actions)
- A3C (Asynchronous Advantage Actor-Critic)
- **TRPO & PPO**
- DDPG (Deep Deterministic Policy Gradient)
 - ▶ TD3
 - ▶ SAC



Trust Region Policy Optimization (TRPO)

- TRPO is a policy optimization algorithm
 - can replace gradient descent
- There are many gradient descent methods
 - Original gradient descent method
 - Natural gradient descent method
 - Stochastic gradient descent method
- TRPO is similar to natural gradient descent method
- TRPO can be combined with A2C, called ACKTR

TRPO

- Consider a Markov decision process (MDP), defined by the tuple

$$(S, A, P, r, \rho_0, \gamma)$$

- S is a finite set of states, A is finite set of actions
 - $P: S \times A \times S \rightarrow \mathbb{R}$ is the transition probability distribution
 - r is reward function
 - $\rho_0: S \rightarrow \mathbb{R}$ is the distribution of initial state (implicitly, $s_0 \sim \rho_0$)
 - $\gamma \in (0, 1)$ is discounted factor
- Let π be a stochastic policy $\pi: S \times A \rightarrow [0, 1]$
- The return function of reinforcement learning is

$$\eta(\pi) := E_{s_0 \sim \rho_0} [V_\pi(s_0)] = \mathbb{E}_{s_0, a_0, \dots \sim \rho_0, \pi} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t) \right]$$



TRPO

- Starting point:

$$\eta(\tilde{\pi}) = \eta(\pi) + \mathbb{E}_{s_0, a_0, \dots \sim \tilde{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t A_{\pi}(s_t, a_t) \right]$$

- Proposed in 2002 by Kakade & Langford
- Note: for simplicity, $\sim \rho_0$ is omitted later.
- This implies that we can derive “return of new policy” from “advantage of old policy”
 - Advantage $A_{\pi}(s_t, a_t) := Q_{\pi}(s_t, a_t) - V_{\pi}(s_t)$

Appendix (Proof of the previous equation)


Since $A_\pi(s, a) = E_{s' \sim P(s'|s, a)}[r(s) + \gamma V_\pi(s') - V_\pi(s)]$, we have

$$\begin{aligned}
& E_{s_0, a_0, \dots \sim \tilde{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t A_{\pi}(s_t, a_t) \right] \\
&= E_{s_0, a_0, \dots \sim \tilde{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t (r(s_t) + \gamma V_{\pi}(s_{t+1}) - V_{\pi}(s_t)) \right] \\
&= E_{s_0, a_0, \dots \sim \tilde{\pi}} \left[-V_{\pi}(s_0) + \sum_{t=0}^{\infty} \gamma^t r(s_t) \right] \\
&= -E_{s_0} [V_{\pi}(s_0)] + E_{s_0, a_0, \dots \sim \tilde{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t) \right] \\
&= -\eta(\pi) + \eta(\tilde{\pi}) \quad \square
\end{aligned}$$



TRPO

- Advantage $A_{\pi}(s_t, a_t) := Q_{\pi}(s_t, a_t) - V_{\pi}(s_t)$
- Can evaluate the current action compared to average value



TRPO

- Expanding $\eta(\tilde{\pi})$, we get

$$\begin{aligned}
 \eta(\tilde{\pi}) &= \eta(\pi) + \mathbb{E}_{s_0, a_0, \dots \sim \tilde{\pi}} \left[\sum_{t=0}^{\infty} \gamma^t A_{\pi}(s_t, a_t) \right] \\
 &= \eta(\pi) + \sum_{t=0}^{\infty} \left(\sum_s \left(P(s_t = s | \tilde{\pi}) \sum_a \tilde{\pi}(a|s) \gamma^t A_{\pi}(s, a) \right) \right) \\
 &= \eta(\pi) + \sum_s \left(\left(\sum_{t=0}^{\infty} \gamma^t P(s_t = s | \tilde{\pi}) \right) \left(\sum_a \tilde{\pi}(a|s) A_{\pi}(s, a) \right) \right) \\
 &\quad \text{Called density of } s, \text{ denoted } \rho_{\tilde{\pi}}(s) \\
 &= \eta(\pi) + \sum_s \left(\rho_{\tilde{\pi}}(s) \left(\sum_a \tilde{\pi}(a|s) A_{\pi}(s, a) \right) \right)
 \end{aligned}$$

- Convert the view from each time point t to each state s



TRPO

$$\rho_{\tilde{\pi}}(s) := \sum_{t=0}^{\infty} \gamma^t P(s_t = s | \tilde{\pi}) = \underbrace{P(s_0 = s) + \gamma P(s_1 = s) + \gamma^2 P(s_2 = s) + \dots}_{\text{unnormalized discounted visitation frequencies}}$$

- Denote the un-normalized discounted visitation frequencies by $\rho_{\tilde{\pi}}(s)$, then the return of $\tilde{\pi}$ become

$$\eta(\tilde{\pi}) = \eta(\pi) + \sum_s \rho_{\tilde{\pi}}(s) \sum_a \tilde{\pi}(a|s) A_{\pi}(s, a)$$

- This implies
 - any policy update $\pi \rightarrow \tilde{\pi}$ that has a nonnegative expected advantage at every state s , is guaranteed to increase the policy performance η
 - or: If all $\sum_a \tilde{\pi}(a|s) A_{\pi}(s, a)$ are non-negative for the new policy $\tilde{\pi}$, the policy performance η must be improved.



TRPO

$$\rho_{\tilde{\pi}}(s) := \sum_{t=0}^{\infty} \gamma^t P(s_t = s | \tilde{\pi}) = \underbrace{P(s_0 = s) + \gamma P(s_1 = s) + \gamma^2 P(s_2 = s) + \dots}_{\text{unnormalized discounted visitation frequencies}}$$

- Denote the un-normalized discounted visitation frequencies by $\rho_{\tilde{\pi}}(s)$, then the return of $\tilde{\pi}$ become

$$\eta(\tilde{\pi}) = \eta(\pi) + \sum_s \rho_{\tilde{\pi}}(s) \sum_a \tilde{\pi}(a|s) A_{\pi}(s, a)$$

- This implies
 - any policy update $\pi \rightarrow \tilde{\pi}$ that has **a nonnegative expected advantage at every state s** , is guaranteed to increase the policy performance η
 - or: If **all $\sum_a \tilde{\pi}(a|s) A_{\pi}(s, a)$ are non-negative for the new policy $\tilde{\pi}$** , the policy performance η must be improved.



TRPO

$$\eta(\tilde{\pi}) = \eta(\pi) + \sum_s \rho_{\tilde{\pi}}(s) \sum_a \tilde{\pi}(a|s) A_{\pi}(s, a)$$

- However, due to the complex dependency of $\rho_{\tilde{\pi}}(s)$ on $\tilde{\pi}$ makes above equation difficult to optimize directly
- Instead, introducing local approximation to η :

$$L_{\pi}(\tilde{\pi}) = \eta(\pi) + \sum_s \rho_{\pi}(s) \sum_a \tilde{\pi}(a|s) A_{\pi}(s, a)$$

- L ignores changes in state visitation density due to changes in the policy
- Key: **try to maximize $L_{\pi}(\tilde{\pi})$ instead of $\eta(\tilde{\pi})$.**
 - Question: why is it fine to replace $\rho_{\tilde{\pi}}(s)$ by $\rho_{\pi}(s)$?



TRPO


- If we have a policy π_θ , which is differentiable w.r.t. θ , then L_π matches η to first order. i.e., for any parameter θ_{old}

$$L_{\pi_{\theta_{old}}}(\pi_{\theta_{old}}) = \eta(\pi_{\theta_{old}}),$$

$$\nabla_\theta L_{\pi_{\theta_{old}}}(\pi_\theta) \Big|_{\theta=\theta_{old}} = \nabla_\theta \eta(\pi_\theta) \Big|_{\theta=\theta_{old}}$$

Proved in next page

- This implies that a step small enough that improves $L_{\pi_{old}}$ will also improve η .



- Sutton's proof by induction for

$$\frac{\partial \eta(\pi_\theta)}{\partial \theta} = \sum_s \rho^\pi(s) \sum_a \frac{\partial \pi_\theta(a|s)}{\partial \theta} Q^\pi(s, a)$$

For the start-state formulation:

$$\begin{aligned} \frac{\partial V^\pi(s)}{\partial \theta} &\stackrel{\text{def}}{=} \frac{\partial}{\partial \theta} \sum_a \pi(s, a) Q^\pi(s, a) \quad \forall s \in \mathcal{S} \\ &= \sum_a \left[\frac{\partial \pi(s, a)}{\partial \theta} Q^\pi(s, a) + \pi(s, a) \frac{\partial}{\partial \theta} Q^\pi(s, a) \right] \\ &= \sum_a \left[\frac{\partial \pi(s, a)}{\partial \theta} Q^\pi(s, a) + \pi(s, a) \frac{\partial}{\partial \theta} \left[\mathcal{R}_s^a + \sum_{s'} \gamma \mathcal{P}_{ss'}^a V^\pi(s') \right] \right] \\ &= \sum_a \left[\frac{\partial \pi(s, a)}{\partial \theta} Q^\pi(s, a) + \pi(s, a) \sum_{s'} \gamma \mathcal{P}_{ss'}^a \frac{\partial}{\partial \theta} V^\pi(s') \right] \quad (7) \\ &= \sum_x \sum_{k=0}^{\infty} \gamma^k P_{\tau}(s \rightarrow x, k, \pi) \sum_a \frac{\partial \pi(x, a)}{\partial \theta} Q^\pi(x, a), \end{aligned}$$



- Sutton's proof by induction for


$$\begin{aligned}\frac{\partial \eta(\pi_\theta)}{\partial \theta} &= \sum_s \rho^\pi(s) \sum_a \frac{\partial \pi_\theta(a|s)}{\partial \theta} Q^\pi(s, a) \\ &= \sum_s \rho^\pi(s) \sum_a \frac{\partial \pi_\theta(a|s)}{\partial \theta} A^\pi(s, a) \\ &\quad (\text{Why? } \sum_a \pi_\theta(a|s) V^\pi(s) = 1) \\ &= \frac{\partial L(\pi_\theta)}{\partial \theta}\end{aligned}$$



TRPO

- And $D_{TV}(p \parallel q)^2 \leq D_{KL}(p \parallel q)$.
- Let $D_{KL}^{max}(\pi, \tilde{\pi}) = \max_s D_{KL}(\pi(\cdot | s) \parallel \tilde{\pi}(\cdot | s))$, then

$$\eta(\tilde{\pi}) \geq L_{\pi}(\tilde{\pi}) - \frac{C \cdot D_{KL}^{max}(\pi, \tilde{\pi})}{4\epsilon\gamma}$$
 where $C = \frac{4\epsilon\gamma}{(1 - \gamma)^2}$
 - When $\pi \rightarrow \tilde{\pi}$, $D_{KL}^{max}(\pi, \tilde{\pi}) \rightarrow 0$, so the lower bound is tight. How much we improve on $L_{\pi}(\tilde{\pi})$, how much the return $\eta(\tilde{\pi})$ also improve
 - When π is not close to $\tilde{\pi}$, the penalty is large since constant C is large, and the lower bound is meaningless.
- A kind of MM algorithm
 - Minorize-Maximization or
 - Majorize-Minimization



TRPO

$$\eta(\tilde{\pi}) \geq L_{\pi}(\tilde{\pi}) - C \cdot D_{KL}^{max}(\pi, \tilde{\pi})$$



- We show that the improvement must be monotonically increasing (MM algorithm)
- Let $M_i(\pi) = L_{\pi_i}(\pi) - C \cdot D_{KL}^{max}(\pi_i, \pi)$:

$$\eta(\pi) \geq M_i(\pi)$$

$$\eta(\pi_i) = M_i(\pi_i)$$

$$\eta(\pi) - \eta(\pi_i) \geq M_i(\pi) - M_i(\pi_i)$$
- Let $\pi_{i+1} = \operatorname{argmax}_{\pi} M_i(\pi)$, then

$$\eta(\pi_{i+1}) - \eta(\pi_i) \geq M_i(\pi_{i+1}) - M_i(\pi_i) \geq 0$$
 and thus the return of next iteration is not worse than current one.



TRPO

● Algorithm

Algorithm 1 Policy iteration algorithm guaranteeing non-decreasing expected return η

Initialize π_0 .

for $i = 0, 1, 2, \dots$ until convergence **do**

 Compute all advantage values $A_{\pi_i}(s, a)$.

 Solve the constrained optimization problem

$$\pi_{i+1} = \arg \max_{\pi} [L_{\pi_i}(\pi) - CD_{\text{KL}}^{\max}(\pi_i, \pi)]$$

$$\text{where } C = 4\epsilon\gamma/(1 - \gamma)^2$$

$$\text{and } L_{\pi_i}(\pi) = \eta(\pi_i) + \sum_s \rho_{\pi_i}(s) \sum_a \pi(a|s) A_{\pi_i}(s, a)$$

end for



TRPO


● Problems:

- In practice, the step size is very small
- D_{KL}^{max} is hard to compute
- How do we approximate the objective function and constraint?

TRPO

- In practice, if using the penalty coefficient C recommended by the theory above, **the step size would be very small.**
- One way to take larger steps in a robust way is to use a constraint on the KL divergence between the new policy and the old policy, i.e., a **trust region constraint**:

$$\begin{aligned} & \max_{\theta} L_{\theta_{old}}(\theta) \\ & \text{subject to } D_{KL}^{max}(\theta_{old}, \theta) \leq \delta \end{aligned}$$



TRPO (can be skipped)


- Since the D_{KL}^{max} is hard to compute, we can use a heuristic approximation which considers the average KL divergence

$$\bar{D}_{KL}^\rho(\theta_{old}, \theta) := \mathbb{E}_{s \sim \rho} \left[D_{KL} \left(\pi_{\theta_{old}}(\cdot | s) \parallel \pi_{\theta}(\cdot | s) \right) \right]$$

- Thus, the problem becomes

$$\max_{\theta} L_{\theta_{old}}(\theta)$$

$$\text{subject to } \bar{D}_{KL}^\rho(\theta_{old}, \theta) \leq \delta$$



TRPO

- Transform the problem: $\max_{\theta} L_{\theta_{old}}(\theta)$

$$\max_{\theta} \sum_s \rho_{\theta_{old}}(s) \sum_a \pi_{\theta}(a|s) A_{\theta_{old}}(s, a)$$

subject to $\bar{D}_{KL}^{\rho}(\theta_{old}, \theta) \leq \delta$

1. Replace $\sum_s \rho_{\theta_{old}}(s)[\dots]$ by expectation $\frac{1}{1-\gamma} \mathbb{E}_{s \sim \rho_{\theta_{old}}}[\dots]$
2. Replace **the sum over the actions** by **an importance sampling estimator**.
Using $\pi_{\theta_{old}}(a|s)$ to denote the sampling distribution, then the contribution of a single s_n to the loss function is:

$$\sum_a \pi_{\theta}(a|s_n) A_{\theta_{old}}(s_n, a) = \mathbb{E}_{a \sim \pi_{\theta_{old}}(a|s_n)} \left[\frac{\pi_{\theta}(a|s_n)}{\pi_{\theta_{old}}(a|s_n)} A_{\theta_{old}}(s_n, a) \right]$$



TRPO

- The problem at the beginning:

$$\begin{aligned} & \max_{\theta} L(\pi_{\theta_{old}}) \text{ or} \\ & \max_{\theta} \sum_s \rho_{\theta_{old}}(s) \sum_a \pi_{\theta}(a|s) A_{\theta_{old}}(s, a) \\ & \text{subject to } \bar{D}_{KL}^{\rho}(\theta_{old}^a, \theta) \leq \delta \end{aligned}$$

- And currently, we solve:

$$\begin{aligned} & \max_{\theta} \mathbb{E}_{s \sim \rho_{\theta_{old}}, a \sim \pi_{\theta_{old}}} \left[\frac{\pi_{\theta}(a|s)}{\pi_{\theta_{old}}(a|s)} A_{\theta_{old}}(s, a) \right] \\ & \text{subject to } \mathbb{E}_{s \sim \rho_{\theta_{old}}} \left[D_{KL}(\pi_{\theta_{old}}(\cdot | s) \parallel \pi_{\theta}(\cdot | s)) \right] \leq \delta \end{aligned}$$

- In another form, maximize a surrogate objective:

$$L^{CPI}(\theta) = \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{old}}(a_t | s_t)} \hat{A}_t \right]$$

- CPI: conservative policy iteration
- \hat{A}_t : can be any form of advantage, like GAE.



Proximal Policy Optimization (PPO)

- Problems of TRPO:
 - still relatively complicated, and
 - not compatible with architectures that include noise (such as dropout) or parameter sharing
- Background:
 - In 2017, OpenAI release a new reinforcement learning algorithms, PPO.
 - PPO has some of the benefits of TRPO, but much simpler to implement, more general, and has better sample complexity.
 - attains the data efficiency and reliable performance of TRPO, while using only first-order optimization
- The experiments show that PPO outperforms other online policy gradient methods, like A2C or TRPO.
 - Although PPO is a little worse than ACER (Actor-Critic with Experience Replay), the implementation of PPO is much easier than ACER.




Generalized Advantage Estimation (GAE)

- Use the learned state-value function $V(s)$ to compute variance-reduced advantage-function estimators.
- PPO uses a truncated version of generalized advantage estimation

$$\hat{A}_t = \delta_t + (\gamma\lambda)\delta_{t+1} + \dots + (\gamma\lambda)^{T-t+1}\delta_{T-1}$$


where $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$



PPO Algorithm

Algorithm 1 PPO, Actor-Critic Style


```
for iteration=1,2,... do  
  for actor=1,2,...,N do  
    Run policy  $\pi_{\theta_{\text{old}}}$  in environment for  $T$  timesteps  
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$   
  end for  
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$   
   $\theta_{\text{old}} \leftarrow \theta$   
end for
```



PPO Algorithm

Algorithm 1 PPO, Actor-Critic Style

```
for iteration=1, 2, ... do
  for actor=1, 2, ..., N do
    Run policy  $\pi_{\theta_{\text{old}}}$  in environment for  $T$  timesteps
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$ 
  end for
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$ 
   $\theta_{\text{old}} \leftarrow \theta$ 
end for
```



Let $\pi_0 = \pi_{\text{old}}$

1. pick a batch with M
2. optimize θ .
from $\pi_i \rightarrow \pi_{i+1}$
3. repeat 1.

$$\pi_0 \rightarrow \pi_1 \rightarrow \pi_2 \rightarrow \dots \rightarrow \pi_K$$



Recall TRPO

- Recall: TRPO maximizes a surrogate objective: $\max_{\theta} L^{CPI}(\theta)$
(with small change on $\pi_{\theta}(a|s)$)

$$L^{CPI}(\theta) = \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)} \hat{A}_t \right] = \hat{\mathbb{E}}_t [r_t(\theta) \hat{A}_t]$$

- CPI: conservative policy iteration
- \hat{A}_t : can be any form of advantage, like GAE.
- Let $r_t(\theta)$ denote the probability ratio (not reward)

$$r_t(\theta) = \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$$

- $r(\theta_{old}) = 1$
- Note: π_{θ} can be any of π_i in PPO

PPO Clip


- Without constraint, the step size of L^{CPI} would be large
- Hence, we consider modifying the objective, to penalize changes to the policy that move $r_t(\theta)$ away from 1
- The main objective proposed in PPO is:

$$L^{CLIP} = \hat{\mathbb{E}}_t [\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t)]$$

- ϵ is a hyper-parameter
- First term implies that the min is L^{CPI}
- Second term modifies the surrogate objective by clipping the probability ratio
- The final objective is a lower bound on L^{CPI}



PPO



● Use one network with same embedding layer:
policy and value

– Value: estimates value of current state by TD-like learning

▶ Value loss: $L_t^{VF}(\theta) = (V_\theta(s_t) - V_t^{target})^2$

– Policy: output probability of actions

▶ Policy obj.: $L_t^{CLIP}(\theta) = \widehat{E}_t [\min(r_t(\theta)\widehat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\widehat{A}_t)]$

where $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$,

\widehat{A}_t is generalized advantage estimation (GAE)


$\widehat{A}_t = \sum_{n=0}^{\infty} (\gamma\lambda)^n \delta_{t+n}^V$,

where $\delta_t^V = r_t + \gamma V_\theta(s_{t+1}) - V_\theta(s_t)$ [TD error]

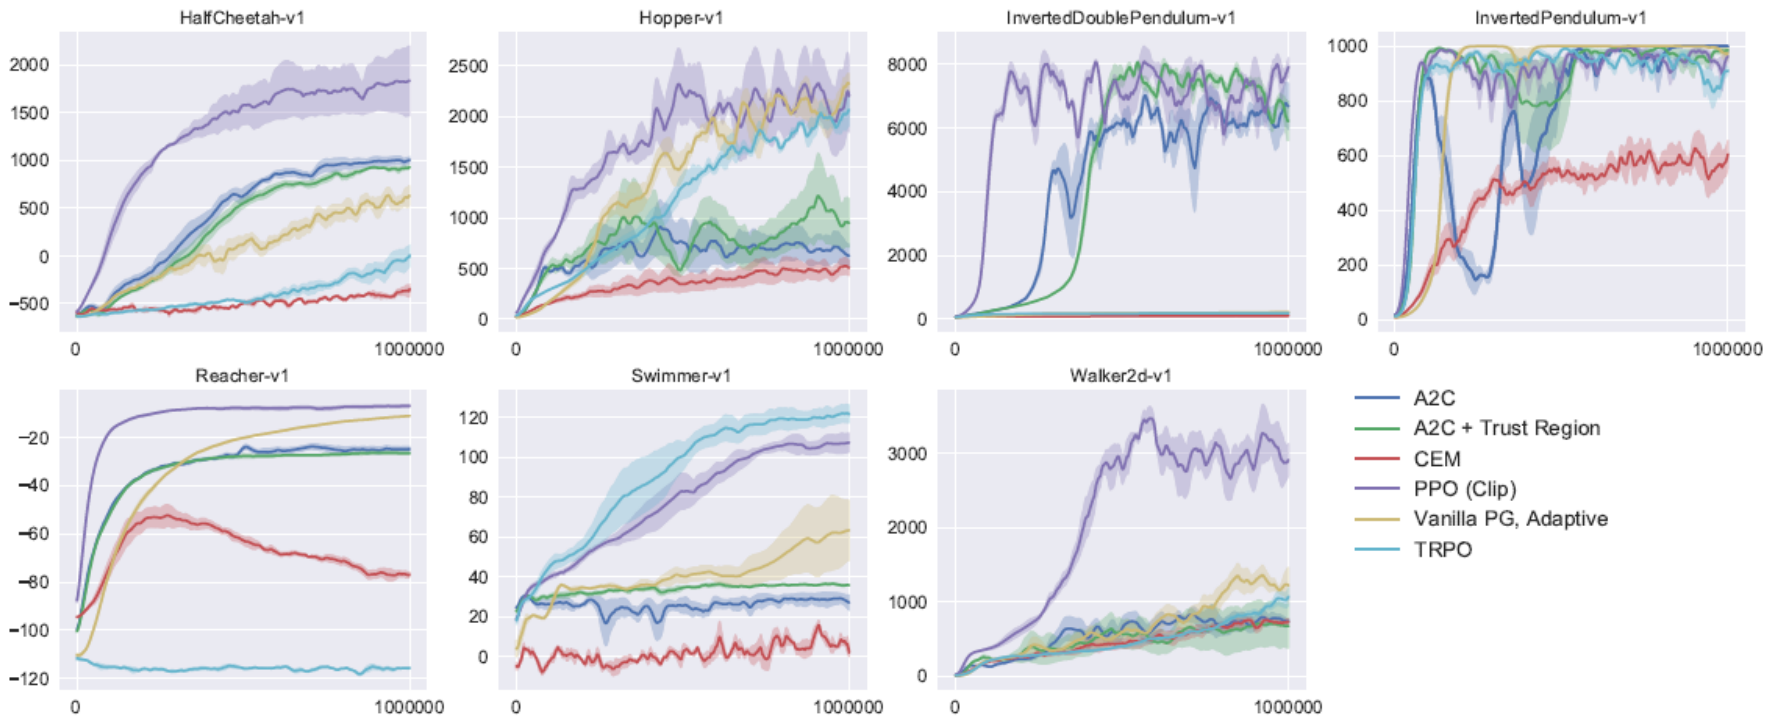
– Total objective (usually version): maximize

$L_t^{CLIP+VF+S}(\theta) = \widehat{E}_t [L_t^{CLIP}(\theta) - c_1 L_t^{VF}(\theta) + c_2 S[\pi_\theta](s_t)]$

▶ Augment with an entropy bonus (S) to ensure sufficient exploration



Experiments - PPO



Policy-Based Reinforcement Learning

- Policy Gradient
- Actor-Critic (Discrete actions)
- A3C (Asynchronous Advantage Actor-Critic)
- TRPO & PPO
- **DDPG (Deep Deterministic Policy Gradient)**
 - ▶ TD3
 - ▶ SAC




Deep Deterministic Policy Gradient (DDPG)

(A Kind of Actor-Critic For Continuous Actions)


- Use two networks: an **actor** and a **critic**
 - **Critic** estimates value of current action by Q-learning

$$\begin{aligned} \nabla_{\omega} L_Q(s_t, a_t | \omega) \\ = \left((r_{t+1} + \gamma Q(s_{t+1}, \mu(s_{t+1} | \theta) | \omega)) - Q(s_t, a_t | \omega) \right) \nabla_{\omega} Q(s_t, a_t | \omega) \end{aligned}$$



- **Actor** updates policy in direction suggested by critic (**DDPG**):

$$\begin{aligned} \nabla_{\theta} J(\mu_{\theta}) &\approx \mathbb{E}_{\mu} [\nabla_{\theta} Q(s_t, \mu(s_t | \theta) | \omega)] \\ &= \mathbb{E}_{\mu} \left[\nabla_a Q(s_t, a | \omega) \Big|_{a=\mu(s_t | \theta)} \nabla_{\theta} \mu(s_t | \theta) \right] \end{aligned}$$



DDPG(1/2)

Behavior and target network

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ
 Initialize **target network** Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$. Initialize replay buffer R
for $t = 1, T$ **do**

Select action $a_t = \mu(s_t|\theta^\mu) + N_t$ **A noise process**

Execute action a_t and observe reward r_t and observe new state s_{t+1}

Experience replay

Store transition (s_t, a_t, r_t, s_{t+1}) in R

Sample random minibatch of M transitions (s_j, a_j, r_j, s_{j+1}) from R

Set $y_i = r_i + \gamma Q'(s_{t+1}, \mu'(s_{t+1}|\theta^{\mu'})|\theta^{Q'})$

Update critic by minimizing the loss: $L = \frac{1}{M} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

Update the actor policy using the sampled gradient:

$$\nabla_{\theta^\mu} \mu|_{s_i} \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$



DDPG(2/2)

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$. Initialize replay buffer R

for $t = 1, T$ **do**

Select action $a_t = \mu(s_t|\theta^\mu) + N_t$

Execute action a_t and observe reward r_t and observe new state s_{t+1}

Store transition (s_t, a_t, r_t, s_{t+1}) in R

Sample random minibatch of M transitions (s_j, a_j, r_j, s_{j+1}) from R

Set $y_i = r_i + \gamma Q'(s_{t+1}, \mu'(s_{t+1}|\theta^{\mu'})|\theta^{Q'})$

**Update the behavior networks
(both actor and critic)**

Update critic by minimizing the loss: $L = \frac{1}{M} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

Update the actor policy using the sampled gradient:

$$\nabla_{\theta^\mu} \mu|_{s_i} \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

Apply “soft” target updates

$$\theta' \leftarrow \tau \theta + (1 - \tau) \theta', \tau \ll 1$$

(0.001 in practice.)


(Note in DQN, θ is copied periodically.

Later, some DQN also used this way)




Experiment Settings

- Run experiments using both a **low-dimensional state description** and **high-dimensional renderings** of the environment
- The frames were downsampled to 64x64 pixels and the 8-bit RGB values were converted to floating point scaled to $[0, 1]$



Example screenshots of a sample of environments to solve with DDPG.

Performance Curves for Those Using Variants of DPG



Light Gray: State Description + Batch Normalization

Dark Gray: State Description + Target Network


Green: State Description + Batch Normalization + Target Network

Blue: Pixels + Target Network





Demo



Policy-Based Reinforcement Learning

- Policy Gradient
- Actor-Critic (Discrete actions)
- A3C (Asynchronous Advantage Actor-Critic)
- TRPO & PPO
- DDPG (Deep Deterministic Policy Gradient)
 - ▶ TD3
 - ▶ SAC



Twin Delayed DDPG (TD3)

Addressing Function Approximation Error in Actor-Critic Methods

Scott Fujimoto, Herke van Hoof and David Meger. “Addressing Function Approximation Error in Actor-Critic Methods.” ICML (2018).

DDPG Overview

initial $\theta, \theta', \phi, \phi'$, replay buffer B

for episode = 1~M **do**

for t = 1~T **do**

Select action using π_ϕ

Play and store transition in B

Sample a batch from B

$$y = r + \gamma Q_{\theta'}(s', \pi_{\phi'}(s'))$$

Update Behavior Critic θ using y

Update Behavior Actor ϕ using **policy gradient**


Update Target

$$\theta' \rightarrow \tau\theta + (1 - \tau)\theta'$$

$$\phi' \rightarrow \tau\phi + (1 - \tau)\phi'$$

	Actor	Critic
Behavior	ϕ	θ
Target	ϕ'	θ'

Network Weight Notation



Method

- Twin Delayed DDPG (TD3)
- TD3 = DDPG + 3 Tricks
 - Clipped Double Q-Learning
 - Delayed Policy Updates
 - Target Policy Smoothing

TD3 Overview

initial $\theta, \theta', \phi, \phi'$, replay buffer B

for episode = 1~M **do**

for t = 1~T **do**

Select action using Critic 1
 θ_1

Play and store transition in B

Sample a batch from B **Trick 1**

$$y = r + \gamma \min_{i=1,2} Q_{\theta'_i}(s', \pi_{\phi'}(s')) + \epsilon \quad \text{Trick 3}$$

Update Behavior Critic θ_1, θ_2 using y

Trick 2 **if** t mod d **then**

Update Behavior Actor ϕ using **policy gradient**



Update Target

$$\theta'_i \rightarrow \tau \theta_i + (1 - \tau) \theta'_i$$

$$\phi' \rightarrow \tau \phi + (1 - \tau) \phi'$$

	Actor	Critic
Behavior	ϕ	θ_1, θ_2
Target	ϕ'	θ'_1, θ'_2

Network Weight Notation



(Recall) Overestimation Problem

- Q-Learning update

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$



Trick 1 : Clipped Double-Q Learning

● Methods to solve overestimation problem

- Double DQN (**Not Good Enough**)

$$y = r + \gamma Q_{\theta'}(s', \pi_{\phi}(s'))$$

- Double-Q Learning (**Not Good Enough**)

$$y_1 = r + \gamma Q_{\theta'_1}(s', \pi_{\phi_1}(s'))$$

$$y_2 = r + \gamma Q_{\theta'_2}(s', \pi_{\phi_2}(s'))$$

	Actor	Critic
Behavior	ϕ	θ_1, θ_2
Target	ϕ'	θ'_1, θ'_2

Network Weight Notation

● Clipped Double-Q Learning (**Better**)

$$\boxed{y} = r + \gamma \min[Q_{\theta'_1}(s', \boxed{\pi_{\phi}(s')}), Q_{\theta'_2}(s', \boxed{\pi_{\phi}(s')})]$$

Only one Q target


Only one actor



Trick 2 : Delayed Policy Updates

- Use **lower frequency** to update **behavior actor** and **target networks**.

```
initial
for episode = 1~M do
  for t = 1~T do
    ...
    Update Behavior Critic
    Update Behavior Actor
    Update Targets Networks
```



```
initial
for episode = 1~M do
  for t = 1~T do
    ...
    Update Behavior Critic
    if t mod d then
      Update Behavior Actor
      Update Targets Networks
```

Hyperparameter



Trick 3 : Target Policy Smoothing

- Assumption
 - Similar actions have similar values

- Add noise to **action value**

$$y = r + \gamma Q(s', \pi(s')) + \epsilon, \epsilon \sim \text{clip}(\mathcal{N}(0, \sigma), -c, c)$$

Hyperparameters

- Regularization

Algorithm 1 TD3

Initialize critic networks $Q_{\theta_1}, Q_{\theta_2}$, and actor network π_ϕ with random parameters θ_1, θ_2, ϕ

Initialize target networks $\theta'_1 \leftarrow \theta_1, \theta'_2 \leftarrow \theta_2, \phi' \leftarrow \phi$

Initialize replay buffer \mathcal{B}

for $t = 1$ **to** T **do**

 Select action with exploration noise $a \sim \pi_\phi(s) + \epsilon$,

$\epsilon \sim \mathcal{N}(0, \sigma)$ and observe reward r and new state s'

 Store transition tuple (s, a, r, s') in \mathcal{B}

 Sample mini-batch of N transitions (s, a, r, s') from \mathcal{B}

$\tilde{a} \leftarrow \pi_{\phi'}(s') + \epsilon, \quad \epsilon \sim \text{clip}(\mathcal{N}(0, \tilde{\sigma}), -c, c)$

$y \leftarrow r + \gamma \min_{i=1,2} Q_{\theta'_i}(s', \tilde{a})$

 Update critics $\theta_i \leftarrow \text{argmin}_{\theta_i} N^{-1} \sum (y - Q_{\theta_i}(s, a))^2$

if $t \bmod d$ **then**

 Update ϕ by the deterministic policy gradient:

$\nabla_\phi J(\phi) = N^{-1} \sum \nabla_a Q_{\theta_1}(s, a)|_{a=\pi_\phi(s)} \nabla_\phi \pi_\phi(s)$

 Update target networks:

$\theta'_i \leftarrow \tau \theta_i + (1 - \tau) \theta'_i$

$\phi' \leftarrow \tau \phi + (1 - \tau) \phi'$

end if

end for

1. Clipped Double Q-Learning for Actor-Critic



Algorithm 1 TD3

Initialize critic networks $Q_{\theta_1}, Q_{\theta_2}$, and actor network π_ϕ with random parameters θ_1, θ_2, ϕ

Initialize target networks $\theta'_1 \leftarrow \theta_1, \theta'_2 \leftarrow \theta_2, \phi' \leftarrow \phi$

Initialize replay buffer \mathcal{B}

for $t = 1$ **to** T **do**

 Select action with exploration noise $a \sim \pi_\phi(s) + \epsilon$,

$\epsilon \sim \mathcal{N}(0, \sigma)$ and observe reward r and new state s'

 Store transition tuple (s, a, r, s') in \mathcal{B}

 Sample mini-batch of N transitions (s, a, r, s') from \mathcal{B}

$\tilde{a} \leftarrow \pi_{\phi'}(s') + \epsilon, \quad \epsilon \sim \text{clip}(\mathcal{N}(0, \tilde{\sigma}), -c, c)$

$y \leftarrow r + \gamma \min_{i=1,2} Q_{\theta'_i}(s', \tilde{a})$

 Update critics $\theta_i \leftarrow \text{argmin}_{\theta_i} N^{-1} \sum (y - Q_{\theta_i}(s, a))^2$

if $t \bmod d$ **then**

 Update ϕ by the deterministic policy gradient:

$\nabla_\phi J(\phi) = N^{-1} \sum \nabla_a Q_{\theta_1}(s, a)|_{a=\pi_\phi(s)} \nabla_\phi \pi_\phi(s)$

 Update target networks:

$\theta'_i \leftarrow \tau \theta_i + (1 - \tau) \theta'_i$

$\phi' \leftarrow \tau \phi + (1 - \tau) \phi'$

end if

end for

1. Clipped Double Q-Learning for Actor-Critic

2. Delayed Policy Updates



Algorithm 1 TD3

Initialize critic networks $Q_{\theta_1}, Q_{\theta_2}$, and actor network π_ϕ with random parameters θ_1, θ_2, ϕ

Initialize target networks $\theta'_1 \leftarrow \theta_1, \theta'_2 \leftarrow \theta_2, \phi' \leftarrow \phi$

Initialize replay buffer \mathcal{B}

for $t = 1$ **to** T **do**

 Select action with exploration noise $a \sim \pi_\phi(s) + \epsilon$,

$\epsilon \sim \mathcal{N}(0, \sigma)$ and observe reward r and new state s'

 Store transition tuple (s, a, r, s') in \mathcal{B}

 Sample mini-batch of N transitions (s, a, r, s') from \mathcal{B}

$\tilde{a} \leftarrow \pi_{\phi'}(s') + \epsilon, \quad \epsilon \sim \text{clip}(\mathcal{N}(0, \tilde{\sigma}), -c, c)$

$y \leftarrow r + \gamma \min_{i=1,2} Q_{\theta'_i}(s', \tilde{a})$

 Update critics $\theta_i \leftarrow \text{argmin}_{\theta_i} N^{-1} \sum (y - Q_{\theta_i}(s, a))^2$

if $t \bmod d$ **then**

 Update ϕ by the deterministic policy gradient:

$\nabla_\phi J(\phi) = N^{-1} \sum \nabla_a Q_{\theta_1}(s, a)|_{a=\pi_\phi(s)} \nabla_\phi \pi_\phi(s)$

 Update target networks:

$\theta'_i \leftarrow \tau \theta_i + (1 - \tau) \theta'_i$

$\phi' \leftarrow \tau \phi + (1 - \tau) \phi'$

end if

end for

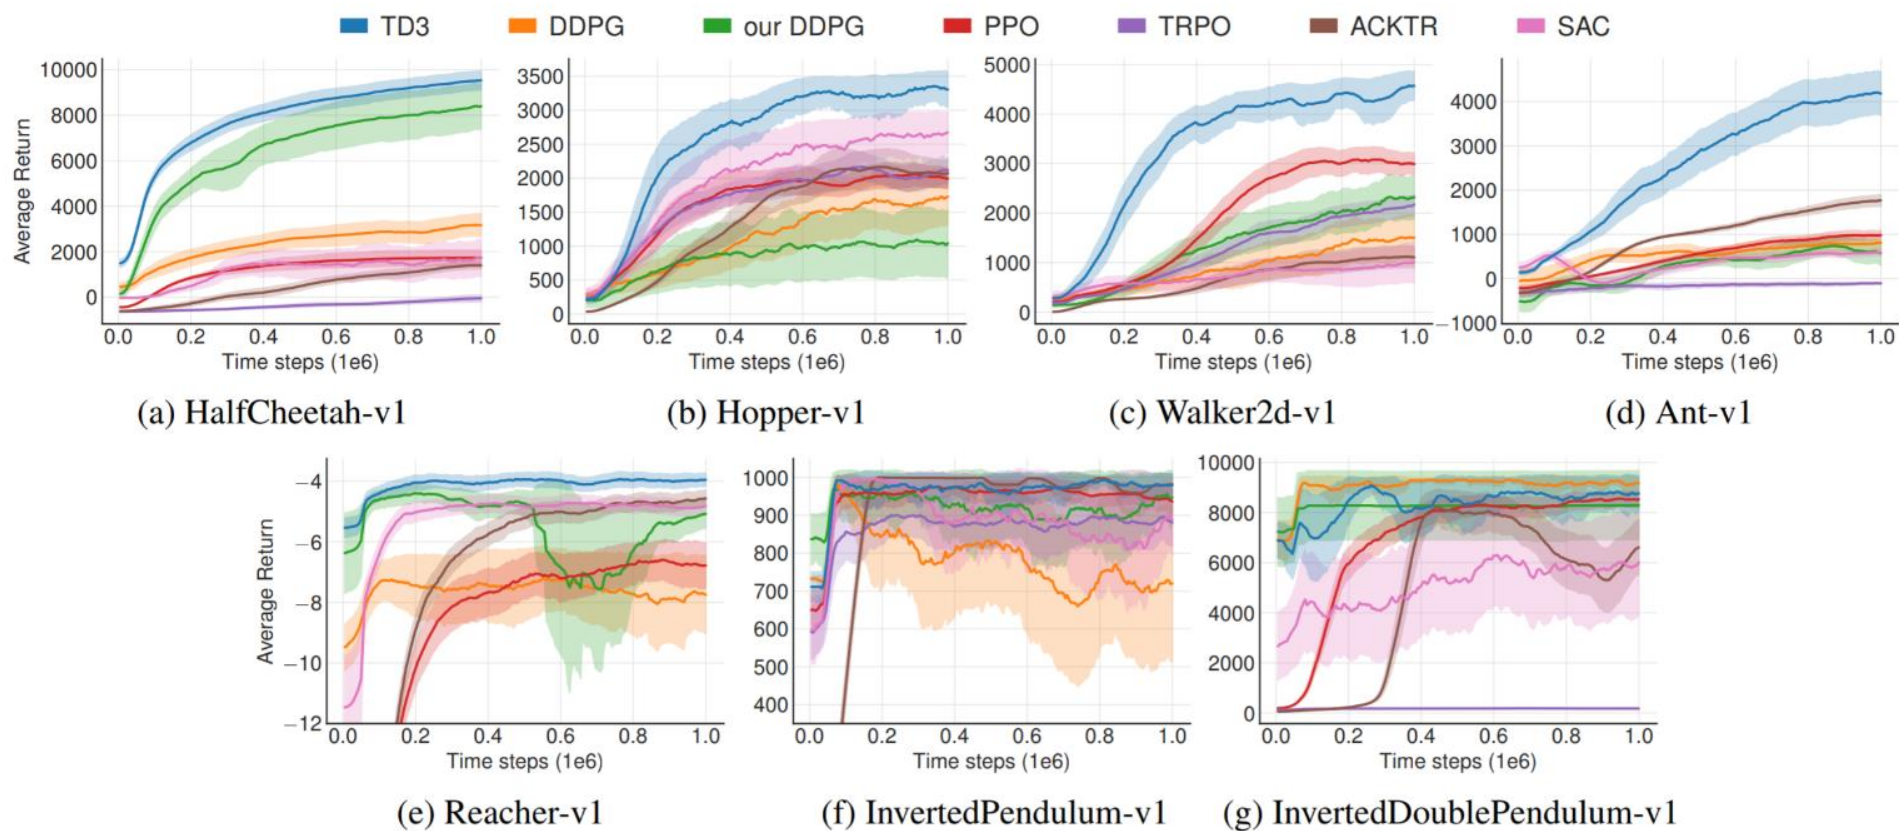
1. Clipped Double Q-Learning for Actor-Critic

2. Delayed Policy Updates

3. Target Policy Smoothing Regularization



Experiment



Experiments: Compared to Others

Environment	TD3	DDPG	Our DDPG	PPO	TRPO	ACKTR	SAC
HalfCheetah	9636.95 \pm 859.065	3305.60	8577.29	1795.43	-15.57	1450.46	2347.19
Hopper	3564.07 \pm 114.74	2020.46	1860.02	2164.70	2471.30	2428.39	2996.66
Walker2d	4682.82 \pm 539.64	1843.85	3098.11	3317.69	2321.47	1216.70	1283.67
Ant	4372.44 \pm 1000.33	1005.30	888.77	1083.20	-75.85	1821.94	655.35
Reacher	-3.60 \pm 0.56	-6.51	-4.01	-6.18	-111.43	-4.26	-4.44
InvPendulum	1000.00 \pm 0.00	1000.00	1000.00	1000.00	985.40	1000.00	1000.00
InvDoublePendulum	9337.47 \pm 14.96	9355.52	8369.95	8977.94	205.85	9081.92	8487.15



Policy-Based Reinforcement Learning

- Policy Gradient
- Actor-Critic (Discrete actions)
- A3C (Asynchronous Advantage Actor-Critic)
- TRPO & PPO
- DDPG (Deep Deterministic Policy Gradient)
 - ▶ TD3
 - ▶ SAC (Soft Actor Critic)



Introduction

- SAC is

- Open-source (by original authors)
 - ▶ <https://sites.google.com/view/sac-and-applications>
- Perform well (as in realistic environment)
- Key idea is easy to understand
 - ▶ Maximum entropy reinforcement learning

Introduction

- Soft actor critic (SAC) train a policy that maximizes a trade-off between expected return and entropy
 - Still getting high performance while acting as random as possible
 - Augment the objective function with entropy term
- Evolution of SAC
 - Soft Q-learning (SQL)
 - Soft Actor-Critic (SAC)
 - Soft Actor-Critic with automating entropy adjustment(SAC)




Problem

- The above methods (PPO, DDPG) focus more on exploitation
 - The objective function is mainly based on the return
 - May be trapped in local optimum **without exploration**

Extremely simple case

Return	Up	Left	Down	Right
	0	10	0	10



Policy	Up	Left	Down	Right
T=0	0.25	0.25	0.25	0.25
T=1	0.2	0.4	0.2	0.2
...				
T=n	0	1	0	0

If we sampled “left” first

Without any exploration,
the chance to sample the “right”
is harder, resulting in the policy
converges to “left” gradually

The agent will be


- either right or left with 100%
- not right and left with 50%



Problem


- Hard exploration case
 - Extend previous “extremely simple case”

The agent will be either right or left for 100%
But not right and left for 50%



Extremely simple case

Hard for agent to discover policy of “right”
May trap in policy of “left”



Hard exploration case

Problem-Solution

- The exploration ability relies on
 - **Random noise** in selected action
- E.g. DDPG
- During training, the action is disturbed with the random noise

Algorithmus 4 : Deep Deterministic Policy-Gradient

Result : policy parameter θ and action-value weights \mathbf{w}

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ and action-value weights $\mathbf{w} \in \mathbb{R}^d$;

Initialize target policy parameter $\theta' \in \mathbb{R}^{d'}$ and target action-value weights $\mathbf{w}' \in \mathbb{R}^d$;

Initialize experience replay memory \mathcal{D} ;

for $episode = 1, M$ **do**

Observe initial state s_0 from environment ;

for $t = 1, T$ **do**

Select action $a_t = \tau(s, \theta_t) + \mathcal{N}_t$;

Observe reward r_t and next state s_{t+1} from environment ;

Store (s_t, a_t, r_t, s_{t+1}) tuple in \mathcal{D} ;

Sample random batch (s_i, a_i, r_i, s_{i+1}) of size B from \mathcal{D} ;

$\delta_i \leftarrow r_i + \gamma \hat{q}(s_{i+1}, \tau(s_{i+1}, \theta'_t), \mathbf{w}'_t) - \hat{q}(s_i, a_i, \mathbf{w}_t)$;

$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \beta \frac{1}{B} \sum_i \delta_i \nabla_{\mathbf{w}} \hat{q}(s_i, a_i, \mathbf{w}_t)$;

$\theta_{t+1} \leftarrow \theta_t + \alpha \frac{1}{B} \sum_i \nabla_{\theta} \hat{q}(s_i, \tau(s_i, \theta_t), \mathbf{w}_t) \nabla_{\theta} \tau(s_i, \theta_t)$;

Update target networks by

$$\theta'_{t+1} \leftarrow v \theta_t + (1 - v) \theta'_t$$

$$\mathbf{w}'_{t+1} \leftarrow v \mathbf{w}_t + (1 - v) \mathbf{w}'_t$$

end

end



Maximum Entropy Reinforcement Learning

- Standard reinforcement learning (RL) objective function:
 - Total expected rewards:

$$J(\pi_\theta) = \sum_t E_{(s_t, a_t) \sim \rho_{\pi_\theta}} [r(s_t, a_t)]$$

where ρ_{π_θ} is data distribution for policy π_θ

- Maximum entropy RL objective function:
 - Augment with entropy term:

$$J(\pi_\theta) = \sum_t E_{(s_t, a_t) \sim \rho_{\pi_\theta}} [r(s_t, a_t) + \alpha H(\pi_\theta(\cdot | s_t))]$$

where α is temperature for importance of the entropy term



Maximum Entropy Reinforcement Learning

- Encourage exploration with entropy term
 - Entropy in loss function: Consider entropy as regularized term
 - ▶ E.g.: PPO


$$L_t^{CLIP+VF+S}(\theta) = \widehat{E}_t[L_t^{CLIP}(\theta) - c_1 L_t^{VF}(\theta) + c_2 S[\pi_\theta](s_t)]$$

The entropy term **only cares the current state**


- Entropy in objective function: Consider entropy as incentivized exploration reward
 - ▶ E.g.: SAC

$$J(\pi_\theta) = \sum_t E_{(s_t, a_t) \sim \rho_{\pi_\theta}} [r(s_t, a_t) + \alpha H(\pi_\theta(\cdot | s_t))]$$

The entropy term **affects following future states by accumulated return**



2a



2b

Soft Q-Learning

- Objective function: Maximum entropy RL

$$J(\pi_\theta) = \sum_t E_{(s_t, a_t) \sim \rho_{\pi_\theta}} [r(s_t, a_t) + \alpha H(\pi_\theta(\cdot | s_t))]$$

- Soft V-function:

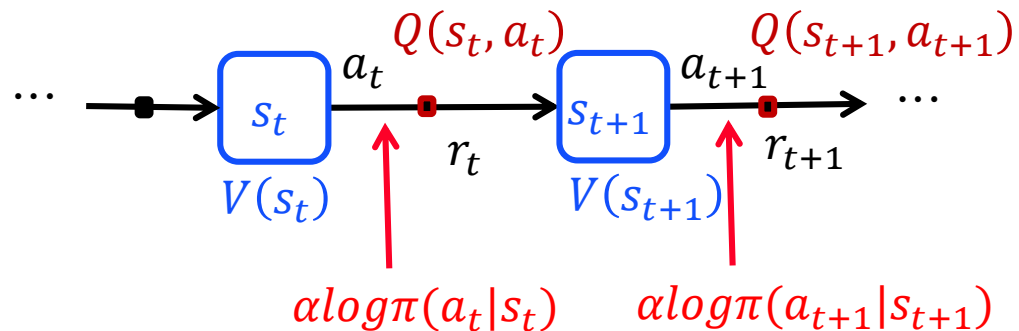
$$V_{soft}(s_t) = E_{a_t \sim \pi_\theta} [Q_{soft}(s_t, a_t) - \alpha \log \pi(a_t | s_t)]$$

- Soft Q-function:

$$Q_{soft}(s_t, a_t) = r_t + \gamma E_{s_{t+1} \sim \rho_{\pi_\theta}} [V_{soft}(s_{t+1})]$$

- Authors prove augment the entropy term still follow Bellman equation property


- Policy evaluation
- Policy improvement
- Policy iteration




Soft Q-Learning

- Gaussian policy:
 - For convenient, usually assume the policy distribution is Gaussian distribution
 - Problem: Not suitable for multimodal case
- Energy-based policy:
 - Use Q value distribution to indicate the policy distribution
 - Assumption: $\pi(a_t|s_t) \propto \exp(Q(s_t, a_t))$

Gaussian policy



Energy-based policy:
Stochastic policy with multimodal



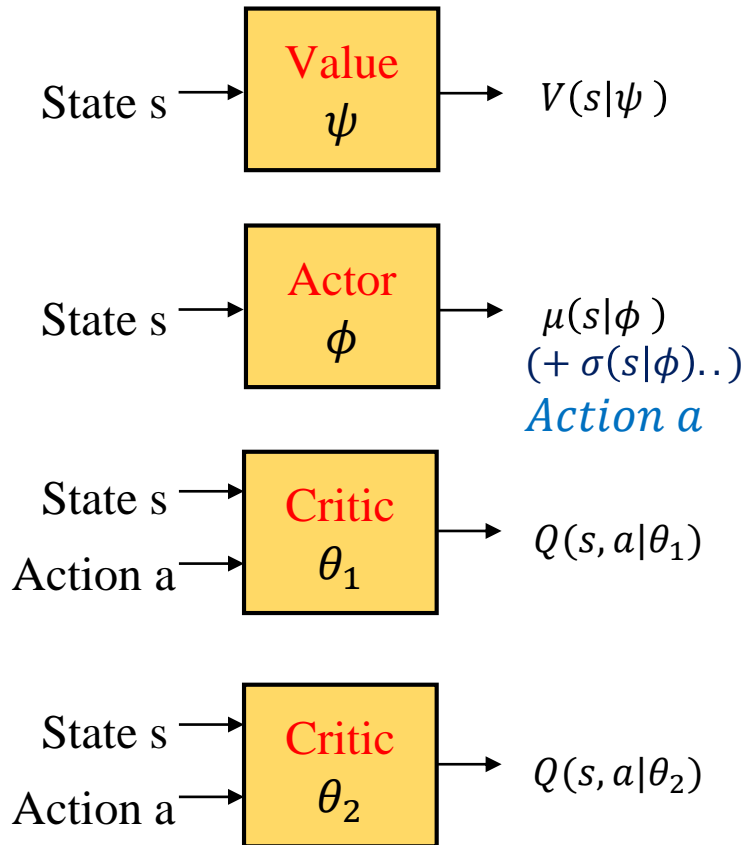
Soft Actor Critic

- Policy: (ideal)


$$\pi(a_t|s_t) = \exp\left(\frac{1}{\alpha}(Q_{soft}(s_t, a_t) - V_{soft}(s_t))\right)$$

- Architecture

- 1 state value (V_ψ) network
- 1 policy network (π_ϕ)
- 2 action-state value (Q-value) network (Q_θ)
 - ▶ Double Q trick: Prevent overestimated in Q
 - ▶ Like TD3



Training of SAC



- D is the distribution of sampled states and actions

- ❑ Value network (V_ψ):

$$J_V(\psi) = E_{s_t \sim D} [\frac{1}{2} (V_\psi(s_t) - \widehat{V}_\psi(s_t))^2]$$

where $\widehat{V}_\psi(s_t) = E_{a_t \sim \pi_\phi}[Q_\theta(s_t, a_t) - \alpha \log \pi_\phi(a_t | s_t)]$


Trained by minimizing the squared residual error (TD error)

- ❑ Q-Value network (Q_θ):


$$J_Q(\theta) = E_{(s_t, a_t) \sim D} \left[\frac{1}{2} \left(Q_\theta(s_t, a_t) - \widehat{Q}_\theta(s_t, a_t) \right)^2 \right]$$

where $\widehat{Q}_\theta(s_t, a_t) = r(s_t, a_t) + \gamma E_{s_{t+1} \sim p}[V_\psi(s_{t+1})]$


Trained by minimizing the **soft Bellman residual error (TD error)**



Training of SAC



- D is the distribution of sampled states and actions
 - Policy network (π_ϕ)
 - Train by **minimizing the KL-divergence**
 - **Use reparameterization trick, sample action from fixed distribution**
- $$J_\pi(\phi) = E_{s_t \sim D, \epsilon_t \sim N} [\log \pi_\phi(f_\phi(\epsilon_t; s_t) | s_t) - Q_\theta(s_t, f_\phi(\epsilon_t; s_t))]$$
- ▶ $a_t = f_\phi(\epsilon_t; s_t)$,
 - ▶ ϵ_t is a noise vector
 - ▶ E.g.: $f_\phi(\epsilon_t; s_t)$ as spherical Gaussian distribution
 - ▶ Take gradient $\nabla_\phi J_\pi(\phi)$



SAC Algorithm

Algorithm 1 Soft Actor-Critic

Initialize parameter vectors $\psi, \bar{\psi}, \theta, \phi$.

for each iteration **do**

for each environment step **do**

$$\mathbf{a}_t \sim \pi_{\phi}(\mathbf{a}_t | \mathbf{s}_t)$$

$$\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$$

$$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$$

end for

for each gradient step **do**

$$\psi \leftarrow \psi - \lambda_V \hat{\nabla}_{\psi} J_V(\psi)$$

$$\theta_i \leftarrow \theta_i - \lambda_Q \hat{\nabla}_{\theta_i} J_Q(\theta_i) \text{ for } i \in \{1, 2\}$$

$$\phi \leftarrow \phi - \lambda_{\pi} \hat{\nabla}_{\phi} J_{\pi}(\phi)$$

$$\bar{\psi} \leftarrow \tau \psi + (1 - \tau) \bar{\psi}$$

end for



end for

Double Q trick




Result


● OpenAI gym v1




(a) Hopper-v1




(b) Walker2d-v1




(c) HalfCheetah-v1



(d) Ant-v1



(e) Humanoid-v1



(f) Humanoid (rllab)



Conclusion

- Soft actor critic (SAC) train a policy that maximize a trade-off between expected return and entropy
 - Still getting high performance while acting as random as possible
- Evolution of SAC
 - Soft Q-learning (SQL)
 - ▶ Soft: $\pi \propto Q(s, a)$
 - Soft Actor-Critic (SAC)
 - ▶ **Argument the objective function with entropy term**
 - Soft Actor-Critic with auto-adjusted temperature (SAC)
 - ▶ **Argument the objective function with entropy term**
 - ▶ Auto-adjust temperature
 - **By constrained policy optimization**

